DEMO STEPS

PART 1:

Terminal 1

1) Start infrastructure
cd ~/Documents/datastream/TP1/project-rag
source ~/Documents/datastream/TP1/datastream/bin/activate
sudo docker compose up -d

2) Clean vector DB (Qdrant)
curl -X DELETE "http://localhost:6333/collections/events"

3) Reset Kafka topics
./reset-topics.sh

To see topics (optional)
sudo docker exec -it kafka kafka-topics --bootstrap-server localhost:9092 --list

4) Start consumers

Terminal 2

cd ~/Documents/datastream/TP1/project-rag
source ~/Documents/datastream/TP1/datastream/bin/activate
python3 index-events.py

Terminal 3

cd ~/Documents/datastream/TP1/project-rag
source ~/Documents/datastream/TP1/datastream/bin/activate
python3 verify-events.py

DURING:

5) Start one ingestion (producer)

Terminal 4

cd ~/Documents/datastream/TP1/project-rag
source ~/Documents/datastream/TP1/datastream/bin/activate
python3 ingest-usgs.py

6) Check Kafka stream

Terminal 5

cd ~/Documents/datastream/TP1/project-rag
source ~/Documents/datastream/TP1/datastream/bin/activate
sudo docker exec -it kafka bash
kafka-console-consumer --bootstrap-server localhost:9092 --topic raw-events --from-beginning --max-messages 10

Terminal 6

cd ~/Documents/datastream/TP1/project-rag
source ~/Documents/datastream/TP1/datastream/bin/activate
sudo docker exec -it kafka bash

kafka-console-consumer --bootstrap-server localhost:9092 --topic verified-events --from-beginning --max-messages 10

Terminal 7
7) Export API key for LLM (before querying)
cd ~/Documents/datastream/TP1/project-rag
source ~/Documents/datastream/TP1/datastream/bin/activate
export GROQ_API_KEY="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"

8) Run RAG query
python3 ask.py "recent earthquakes in California" --minutes 60 --topk 5 --llm
python3 ask.py "recent earthquakes" --topk 5 --llm

FOR THE LATENCY:
/usr/bin/time -f "Latency: %e seconds" python3 ask.py "recent earthquakes in California" --minutes 60 --topk 5 --llm

New terminal:
cd ~/Documents/datastream/TP1/project-rag
source ~/Documents/datastream/TP1/datastream/bin/activate

PART 2: Testing commit-after-write

1. Watch lag (group = index-events):

sudo docker exec -it kafka kafka-consumer-groups --bootstrap-server localhost:9092 --describe --group index-events

GROUP TOPIC PARTITION CURRENT-OFFSET LOG-END-OFFSET LAG CONSUMER-ID HOST CLIENT-ID

index-events verified-events 0 15 15 0 kafka-python-2.2.15-a2a17394-9847-404e-9d8a-c04e81c1369e /172.18.0.1 kafka-python-2

2. Stop Qdrant (cause upsert failure):

sudo docker stop qdrant

To confirm it is stopped:
docker ps | grep qdrant

3. Generate new messages (so verified-events grows):

python3 ingest-eonet.py

(verify-events should pass them through into verified-events)

Lag should increase / not drain while Qdrant is down.

4. Run again:

sudo docker exec -it kafka kafka-consumer-groups --bootstrap-server localhost:9092 --describe --group index-events

5. Restart Qdrant:

sudo docker start qdrant

6. After a few min run again:

sudo docker exec -it kafka kafka-consumer-groups --bootstrap-server localhost:9092 --describe --group index-events

luciana@luciana-ZenBook-UX434DA-UM433DA:~/Documents/datastream/TP1/project-rag$ sudo docker exec -it kafka kafka-consumer-groups --bootstrap-server localhost:9092 --describe --group index-events

GROUP TOPIC PARTITION CURRENT-OFFSET LOG-END-OFFSET **LAG** CONSUMER-ID **HOST CLIENT-ID**
index-events verified-events 0 15 15 0
kafka-python-2.2.15-a2a17394-9847-404e-9d8a-c04e81c1369e /172.18.0.1 kafka-python-2.2.15
luciana@luciana-ZenBook-UX434DA-UM433DA:~/Documents/datastream/TP1/project-rag$ sudo docker exec -it kafka kafka-consumer-groups --bootstrap-server localhost:9092 --describe --group index-events

(new ingest running)

GROUP TOPIC PARTITION CURRENT-OFFSET LOG-END-OFFSET **LAG** CONSUMER-ID **HOST CLIENT-ID**
index-events verified-events 0 15 279 264
kafka-python-2.2.15-a2a17394-9847-404e-9d8a-c04e81c1369e /172.18.0.1 kafka-python-2.2.15
luciana@luciana-ZenBook-UX434DA-UM433DA:~/Documents/datastream/TP1/project-rag$

(a few min after restarting qrant)

GROUP TOPIC PARTITION CURRENT-OFFSET LOG-END-OFFSET LAG CONSUMER-ID HOST CLIENT-ID
index-events verified-events 0 282 282 **0**
kafka-python-2.2.15-a2a17394-9847-404e-9d8a-c04e81c1369e /172.18.0.1 kafka-python-2.2.15